

Adversarial Text Generation for Social Bots

Haris Habibullah

Computational Engineering
haris.habibullah@fau.de

Qiang Ma

Computational Engineering
qiang.ma@fau.de

Jiarong Pan

Informations-und-Kommunikationstechnik
gary.pan@fau.de

Vishal Sukumar

Computational Engineering
vishal.sukumar@fau.de

Abstract—In recent years, social bot has become an important means of dissemination of fake news, malicious speech, and manipulation of public opinion. It becomes more and more important to tackle this problem and analysing the textual pattern of social bot is significantly important for this. In this report, we present that how to train generative adversarial network (GAN) using sequence-GAN (seqGAN) to generate social bot text and demonstrate the result of the network. Experiment results on the chosen dataset demonstrate that this model consistently shows significant improvement over the baseline and a high quality and accuracy of generating texts.

Index Terms—SeqGAN, social bots, twitter, text generation

I. INTRODUCTION

Social media becomes part of our regular life. We spend plenty of time in reading and accepting the information, meanwhile, those media platforms such as Facebook, Twitter are filled with fake, useless or repetitive information. The twitter bot is a type of bot software that controls Twitter accounts via Twitter APIs, which can autonomously perform actions such as tweeting, re-tweeting, and/or sending message to other accounts. To distinguish such bots, it is important to understand their working principle and tweet generation mechanism. The need for detecting the social bots on the popular media is becoming increasingly important. Previous researches are focusing on two directions: how to generate text and find the different behaviour between the real and bot accounts. Applying SeqGAN to generate text has proved to be a success. In this report, we try to use the dataset which collects the tweets from both two kinds of accounts, feeding the data to our model in order to combine those two directions.

We are focusing on the twitter and generating real-like tweets, which can fool a discriminator. After implementing and improving the basic function using SeqGAN, it will help researchers to mimic the basic working of a social bot then learn how to detect them.

Extensive experiments based on synthetic and real data are conducted to investigate the efficacy and properties of the proposed SeqGAN. In our synthetic data environment, SeqGAN significantly outperforms the maximum likelihood methods, scheduled sampling. In real world tasks, i.e. twitter generation, speech language generation and music generation, SeqGAN significantly outperforms the compared baselines in various metrics including human expert judgement.

II. GENERATIVE ADVERSARIAL NETWORK

Generative adversarial network, proposed by [1] belongs to the set of generative models. Figure 1 shows the basic




Fig. 1. Both the generator and the discriminator are neural networks. The generator output is connected directly to the discriminator input. Through backpropagation, the discriminator's classification provides a signal to the generator which is used to update its weights.

framework of GAN. This network has two parts:

- The generator learns to generate plausible data. The generated instances become negative training examples for the discriminator.
- The discriminator learns to distinguish the generator's fake data from real data. The discriminator penalizes the generator for producing implausible results.

When training begins, the generator produces obviously fake data, and the discriminator quickly learns to tell that it's fake. The discriminator in a GAN is simply a classifier. It tries to distinguish real data from the data created by the generator. The discriminator's training data comes from two sources:

- **Real data** instances, such as real pictures of people. The discriminator uses these instances as positive examples during training.
- **Fake data** instances created by the generator. The discriminator uses these instances as negative examples during training.

The generator tries to minimize the following function while the discriminator tries to maximize it:

$$\mathbb{E}_x[\log(D(x))] + \mathbb{E}_z[\log(1 - D(G(z)))] \quad (1)$$

In this function:

- $D(x)$ is the discriminator's estimate of the probability that real data instance x is real.
- \mathbb{E}_x is the expected value over all real data instances.
- $G(z)$ is the generator's output when given noise z .
- $D(G(z))$ is the discriminator's estimate of the probability that a fake instance is real.




Fig. 2. The illustration of SeqGAN. The left part shows that the Discriminator D is trained over the real data and the generated fake data by generator G . The right part shows that the generator G is trained by policy gradient where the final reward is provided by discriminator D and is passed back to the intermediate action value via Monte Carlo search.[?]

- \mathbb{E}_x is the expected value over all random inputs to the generator (in effect, the expected value over all generated fake instances $G(z)$).
- The formula derives from the cross-entropy between the real and generated distributions.

The generator can't directly affect the $\log(D(x))$ term in the function, so for the generator, minimizing the loss is equivalent to minimizing $\log(1 - D(G(z)))$.

III. SEQGAN

In this project, we have used the SeqGAN architecture by ? [?]. It generates text contents similar to the true data by using generative adversarial network, which is tweet from social bots in our case.

There are two main reasons for using this architecture instead of the baseline GAN architecture. First, GAN is designed for generating real-valued, continuous data but has difficulties in directly generating sequence of discrete tokens. Secondly, GAN can only give the score/loss for an entire sequence when it has been generated. It is important to decide if the generated sequence is good enough for now or we need to keep generating. Thus, GAN is not efficient for partial-sequence scoring.

A. Overall Framework

Figure 2 illustrates the basic framework of SeqGAN. The generative model is treated as an agent of reinforcement learning, the state is the generated tokens so far and the action is the next token to be generated. The discriminator D is to evaluate the sequence and feedback the score/loss to guide the generative model. Monte Carlo search is employed to approximate the state-action value.

The sequence generation problem can be formalized as the following. Given a dataset of real-world structured sequences, train a θ -parameterized generative model G_θ to produce a sequence $Y_{1:T} = (y_1, \dots, y_t, \dots, y_T)$, $y_t \in \mathcal{Y}$, where \mathcal{Y} is the vocabulary of candidate tokens. In timestep t , the state s is the current produced tokens (y_1, \dots, y_{t-1}) and the action a is the next token y_t to select.

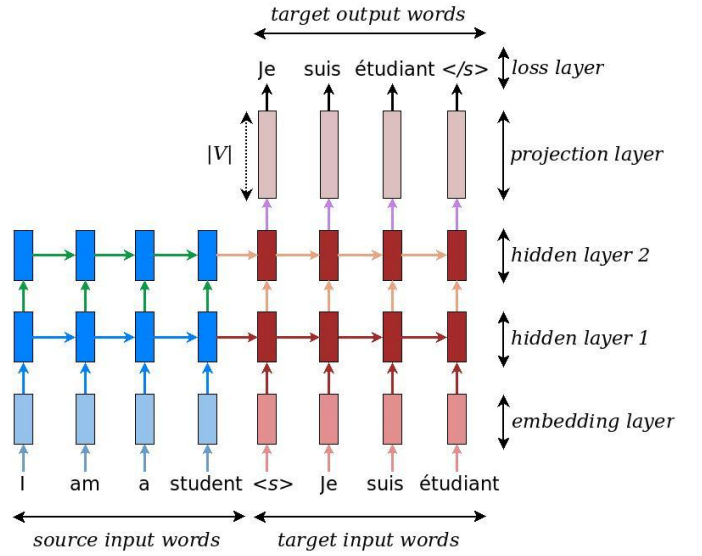


Fig. 3. The illustration of Seq2Seq model by [?]. In our case, in order to generate random text. The source input words would be random cell states and hidden states. The projection layer corresponds to the softmax output layer which maps our hidden states into the output token distribution.

The objective of the generator model $G_\theta(y_t|Y_{1:t-1})$ is to generate a sequence from the start state s_0 to maximize its expected end reward:

$$J_\theta = \mathbb{E}[R_T|s_0, \theta] = \sum_{y_1 \in \mathcal{Y}} G_\theta(y_t|s_0) \cdot Q_{D_\phi}^{G_\theta}(s_0, y_1) \quad (2)$$

where R_T is the reward for a complete sequence. The reward is from the ϕ -parameterized discriminator D_ϕ . $Q_{D_\phi}^{G_\theta}(s_0, y_1)$ is the action-value function of a sequence, i.e. the expected accumulative reward starting from state s , taking action a , and then following policy G_θ . Note that, the reward is considered as the estimated probability of being real by the Discriminator. This objective function implies that starting from a given initial state, the goal of the Generator is to generate a sequence which would make the discriminator consider it is real.

B. The Generative Model for Sequences

The SeqGAN use recurrent neural networks (RNNs) (? [?]) as the generative model. An RNN maps the input embedding representations x_1, \dots, x_T of the sequence x_1, \dots, x_T into a sequence of hidden states h_1, \dots, h_T by using the update function g recursively.

$$h_T = g(h_{t-1}, x_T) \quad (3)$$

Moreover, a softmax output layer maps the hidden states into the output token distribution. In our implementation, we use seq2seq model by ? [?] as our RNN. Figure 3 shows the basic architecture of a seq2seq model.

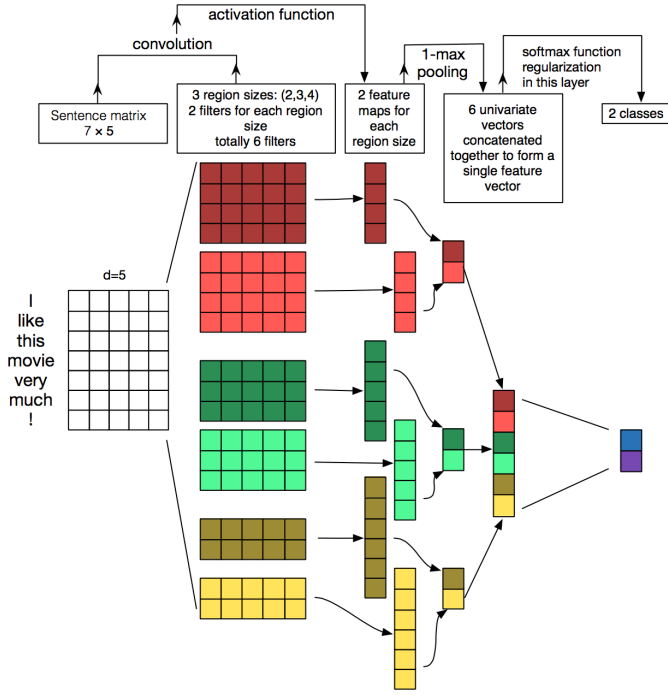


Fig. 4. This is an example of how a CNN classifier work. We concatenate the token embedding vector into a 2 dimensional matrix. Then we perform convolution by various number of kernels with different window size over the matrix to extract different features. After that, we apply a max-over-time pooling operation over the feature map. Finally, a fully connected layer with sigmoid activation is used to output the probability that the input sequence is real.

C. The Discriminative Model for Sequences

In this project, CNN is chosen as our discriminator as CNN has recently been shown of great effectiveness in text (token sequence classification). The optimization target is to minimize the cross entropy between the ground truth label and the predicted probability. Figure 4 demonstrate a simple example of a CNN classifier on text.

IV. RELATED WORK

Following the limitation of GAN, pertinent to the generation of discrete data, such as words, various methods have been proposed to generate a sequence. One of the most popular methods is of the Recurrent Neural Network (RNN), which is used to produce a sequence of tokens and maximize their likelihood given a context [?] [?] .

[?] [?] used the cross-entropy loss function to train the model using Maximum Likelihood Estimation (MLE). The process starts by feeding in words of a given sentence at each epoch and the expected output for each step becomes serves as the next word in the sentence.

[?] and [?] used a similar approach. MLE was used to train the generator in the pre-train stage, while the cross-entropy loss was used to train the discriminator in the pre-train stage. The benefit that the cross-entropy loss function provides is that as the difference between the true label and the predicted label (in terms of probability) increases, the discriminator loss

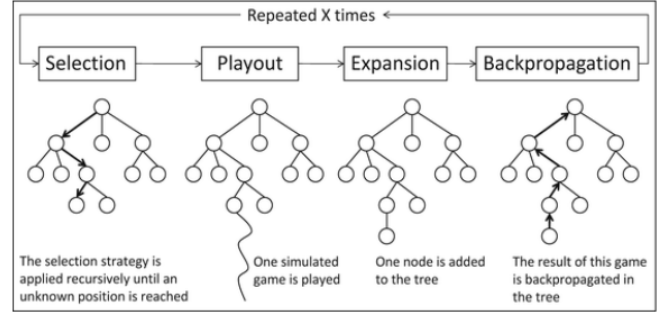


Fig. 5. MCTS working algorithm

increases as well. In this way, an objective for the loss function can be designed, which can minimize the cross-entropy loss.

$$\min_{\phi} -\mathbb{E}_{Y \sim p_{data}} [\log D_{\phi}(Y)] - \mathbb{E}_{Y \sim G_{\theta}} [\log(1 - D_{\phi}(Y))] \quad (4)$$

A number of researchers have worked on devising a reward function for the generator. The challenge here is that during the early stages of generation, a sequence of words will almost always be treated as a bad sequence, discrediting the information in the sequence. For this reason, step-wise reward algorithms have been proposed. [?] [?] used the step-wise GAN (stepGAN) to reward the generator after every step. A single epoch is divided into multiple steps, and evaluation is done based on results / word sequence generated in the particular step with the step input.

A similar approach is carried out in the Monte Carlo Tree Search (MCTS) algorithm. MCTS generate random roll outs, and a reward is generated for each word [?] [?] , as shown in figure 5. Word probabilities at the end of a state are back-propagated to the selection stage, and the word with the highest probability is selected as the next word in the sequence.

V. EXPERIMENT

A. Experiment Design

Our study has used already labelled data, available online from the Social Honeypot Dataset¹ by [?] [?] . The dataset comprised of more than 3 million tweets from legitimate (real) users and approximately 2 million bot tweets. The real tweet data was reduced to approximately 100k tweets for computational ease.

The Experiment was conducted in two stages; Pre-Training and Training. For the Pre-Training phase, an vocab file was created for the tweets using the One-Hot Vector method. This allowed both the discriminator and generator to learn English words and produce similar words or combination of words and not just garbage text.

B. Training Algorithm

Algorithm 1 shows full details about the SeqGAN, referenced from [?] . At the beginning of the training, we use

¹<http://infolab.tamu.edu/data/>

Algorithm 1: Sequence Generative Adversarial Nets

Require : generator policy G_θ ; rollout policy G_β ;
discriminator D_ϕ ; a sequence dataset $S = \{X_{1:T}\}$
Initialize G_θ, G_β with random weights θ, ϕ .
Pre-Train G_θ using MLE on S
 $\beta \leftarrow \theta$
Generate negative samples using G_θ for training D_ϕ
Pre-train D_ϕ via minimizing cross entropy
repeat
for g_steps **do**
Generate a sequence $Y_{1:T} = (y_1, \dots, y_T) \sim G_\theta$
for t in $1:T$ **do**
| Compute $Q(a = y_t; s = Y_{1:t-1})$
end
Update Generator parameters via Policy Gradient
end
for d_steps **do**
Use current G_θ to generate negative samples and
combine with the given positive samples S
Train Discriminator D_ϕ for k epochs by Eq. (4)
end
 $\beta \leftarrow \theta$
until SeqGAN converges

the Maximum Likelihood Estimation (MLE) to pre-train the Generator G_ϕ on the dataset S . After the pre-training the generator and the discriminator are trained alternatively. As the generator gets progressed on g -steps updates, the discriminator has to be trained periodically to be able to identify the real and fake samples effectively. When training the discriminator the positive samples are from the dataset and the negative samples are obtained from the generator. In order to maintain balance we use equal number of positive and negative samples to train the discriminator. The reward function generates Intermediate Rewards, i.e Token wise rewards for partial sequences. Use of MCTS

C. Results

Initially, we pre-trained with the entire dataset, hoping to improve the efficiency during the adversarial training stage. However, this resulted in the generator learning the entire dataset and mimicking the input tweets rather than generating new sequences.

We then used 20% dataset for the pre-training and 80% for the adversarial training. This improved both the processing time and the discriminator ability to distinguish between real and generated-tweets, thus, improving the overall performance of the Generator also.

Figure 6 presents the training results. It can be seen that during the initial epochs, the training losses for both the Generator and the Discriminator fluctuates considerably around 0.5 (the midpoint). This indicates the learning battle between the Generator and the Discriminator, where each tries to outperform the other. After 120 epochs, both the Discriminator and Generator losses are closed to 0.5. This indicates a 50%




Fig. 6. Training loss curves for Discriminator and Generator

probability for a Discriminator to catch a fake, i.e. computer-generated tweet, while a 50% probability for the Generator to fool the Discriminator and get labelled as a legitimate tweet. The Generator loss at the end of the training is still away from the ideal mark, because of the presence of rare and unknown words in the tweets, which are hard to learn for the Generator.

Some of the sample tweets generated by the trained generator are presented in Table I. Though, some tweets may not make a complete sense logically to humans, but the majority of the tweets generated were grammatically correct, and it would be very hard for a bot-detector to classify them as computer-generated tweets. This further stresses on the strength of the adversarial network to learn word sequencing, which more often than not also offers a semantically correct presentation.

TABLE I
TWEETS GENERATED BY THE GENERATOR AFTER 100 EPOCHS

Generated Tweets	Related Real Tweets
Ending slavery, Betty White House.	One group wanted to abolish slavery while other was talking about states rights & "property".
Put with Donald Trump for the future leaks in Philadelphia.	Donald Trump just fired four in one fell swoop in The Apprentice USA! Makes Suralan look meek and mild!
We did n't have transgender bathroom ever to soda after massive much	Internet Radio kicks off Transgender Awareness Week.

VI. CONCLUSION

We were able to create an Adversarial Network that can learn from text and create new text sequences, which adhere to the English grammar parameters. One of the limitations of the network is that the current algorithm is dependent on the initial parameters. If we pre-train with a large (actual) dataset, the generator memorizes everything and creates the same texts rather than using the word-embedding technique. This was taken care by dividing the tweet dataset into proportionate chunks for pre-training and adversarial training.

Furthermore, in the final epoch, the Generator still produces "UNK" labels for the unknown/ rare words. One technique to remove this is of manually replacing all "UNK" labels with related words using the embedded maps. A selection algorithm can be used to provide this functionality.

VII. ACKNOWLEDGEMENTS

We sincerely thank Prof. Dr. Stefan Evert, Dr. Vincent Christlein, Anguelos Nicolaou, Ronak Vijaypal Kosti, and Prof. Dr.-Ing. Andreas Maier for many helpful discussions and comments on the project and the manuscript.